docker Documentation

Release 6.1.0.dev0

The sean developers

September 15, 2015

Contents

Docker Basic

written by sean

Github |

Chapter 0: About

Linux is the best software for software paradise

1.1 Thanks to

- sean
- Mr Ju SS
- OSS Members

1.2 SEAN's Paradise

I think that My Life as Software Engineer was torrible , but it's role for social is important so, I keep going for better life & software development

chapter 1: Docker

2.1 1.1 Basic

2.1.1 1.1.1 Linux

Automatic Install Script

```
$ sudo wget -qO- https://get.docker.com/ | sh
```

remove hell-world

```
$ sudo docker rm `sudo docker ps -aq`
$ sudo docker rmi hello-world
```

.

Ubuntu

Manual install for Ubuntu4.04

```
$ sudo apt-get update
$ sudo apt-get install docker.io
$ sudo ln -sf /usr/bin/docker.io /usr/local/bin/docker
```

RedHat Enterprise Linux, CentOS

CentOS 6

```
$ sudo yum install http://dl.fedoraproject.org/pub/epel/6/x86_64/epel-release-6-8.noarch.rpm
$ sudo yum install docker-io
```

CentOS 7

```
$ sudo yum install docker
```

Docker service execution in CentOS 6.5

```
$ sudo service docker start
```

auto execution during boot in CentOS 6.5

```
$ sudo chkconfig docker on
```

Docker service execution in CentOS 7

```
$ sudo systemctl list-unit-files --type=service |grep docker
$ sudo systemctl enable docker.service
$ sudo systemctl start docker.service
$ sudo systemctl status docker.service
```

.

2.1.2 1.1.2 Mac OS X

https://github.com/boot2docker/osx-installer/releases13 Boot2Docker-1.x.x.pkg

2.1.3 1.1.3 windows

https://github.com/boot2docker/windows-installer/releases52

docker-install.exe

2.2 1.2 Installation

2.2.1 1.2.1 docker default directory

.

- docker default directory change

will create in /var/lib/docker

In CentOS 6.5

```
service docker stop
mkdir /data/docker   (new directory)
vi /etc/sysconfig/docker
```

add following line

```
other_args=" -g /data/docker -p /var/run/docker.pid"
other_args=" -g /docker/data -p /var/run/docker.pid"
```

then save the file and start docker again

```
service docker start
```

and will make repository file in /data/docker

In CentOS 7.0

```
systemctl stop docker.service
vi /etc/sysconfig/docker
```

add following line

```
OPTIONS='-g /docker/data -p /var/run/docker.pid'
```

. and service restart

```
systemctl start docker.servce
```

.

2.2.2 1.2.2 Kernel Upgrade 2.6->3.8

```
yum install http://www.elrepo.org/elrepo-release-6-5.el6.elrepo.noarch.rpm
yum --enablerepo=elrepo-kernel install kernel-ml
```

.when remote access

> cannot access if kernel is not upgrade

*KVM issue

(1) As of kernel-ml-3.10.5-1.el6.elrepo, kernel-ml installed as a KVM guest will panic upon booting (FATAL: Module scsi_wait_scan not found error). This is because virtio_blk is not in the initramfs. More details can be found in:

http://elrepo.org/bugs/view.php?id=401 (external link) https://bugzilla.kernel.org/show_bug.cgi?id=60758 (external link)

A workaround is to rebuild initramfs with a "–add-drivers virtio_blk" option. For example:

dracut –add-drivers virtio_blk -f /boot/initramfs-3.10.5-1.el6.elrepo.x86_64.img 3.10.5-1.el6.elrepo.x86_64

dracut –add-drivers virtio_blk -f /boot/initramfs-4.0.0-1.el6.elrepo.x86_64.img 4.0.0-1.el6.elrepo.x86_64

dracut –add-drivers virtio_blk -f /boot/initramfs-3.19.1-1.el6.elrepo.x86_64.img 3.19.1-1.el6.elrepo.x86_64

dracut –add-drivers virtio_blk -f /boot/initramfs-3.10.71-1.el6.elrepo.x86_64.img 3.10.71-1.el6.elrepo.x86_64 dracut –add-drivers virtio_blk -f /boot/initramfs-4.1.5-1.el6.elrepo.x86_64.img 4.1.5-1.el6.elrepo.x86_64

*cannot found ko.map XXXX cannot resolve

vi /boot/grub/grub.conf

KEYTABLE=ko ==> KEYTABLE=us

- zsh

yum list kernel* xxx yum shell >list kernel*

2.2.3 1.2.3 docker start error

```
usr/bin/docker: relocation error: /usr/bin/docker: symbol dm_task_get_info_with_deferred_remove,
version Base not defined in file libdevmapper.so.1.02 with link time reference
```

.

```
yum-config-manager --enable public_ol6_latest

yum install device-mapper-event-libs
```

.

2.2.4 1.2.4 Build your own image from CentOS

```
yum install feboostrap
febootstrap -i iputils -i vim-minimal -i iproute -i bash -i coreutils -i
yum centos centos http://centos.mirror.iweb.ca/6.4/os/x86_64/ -u http://centos.mirror.iweb.ca/6.4/upd
```

and

```
[root@banshee ~]# cd centos/
[root@banshee centos]# tar -c . | docker import - centos
```

or ISO mount

```
# mkdir rootfs
# mount -o loop /path/to/iso rootfs
# tar -C rootfs -c . | docker import - rich/mybase
```

using osirrox

```
yum install xorriso
osirrox -indev blahblah.iso -extract / /tmp/blahblah
tar -C /tmp/blahblah -cf- . | docker import blahblah
```

- save docker images to tar

```
docker save ubuntu > /tmp/ubuntu.tar
```

extract ubuntu.tar and jump to lagest directory and will see layer.tar

- tar to docker image import

```
cat exampleimage.tgz | docker import - exampleimagelocal:new
```

.

2.2.5 1.2.5 docker images delete

*(none) image delete

```
$ docker rmi $(docker images -f dangling=true | awk '{ print $3 }' | grep -v IMAGE)
```

*all container delete

```
$ sudo docker rm $(docker ps -a -q)
```

*all image delete

```
$ sudo docker rmi -f $(docker images -q)
```

.

2.2.6 1.2.6 gunicorn error

Next we need to install gunicorn. for this we have (as always) several choices.

1. Using YUM. I personally don't recommend it. I know some are happy to use the system packaging management wherever possible, but as for python I don't think it's the way to go.

To install gunicorn using yum:

```
yum install python-gunicorn
```

2. Using easy_install. using easy_install is a better choice for my taste to install python packages. this is how you install gunicorn using easy_install, but I recommend installing gunicorn using PIP as I will show next. . .

```
yum install python-setuptools
easy_install gunicorn
```

3. Using PIP: This is my RECOMMENDED way of installing gunicorn. to install PIP you actually need easy_install so the commands are:

```
yum install python-setuptools
easy_install pip
pip install gunicorn
```

.

- intall from source

```
yum erase python-pip
yum install xz-libs
```

Let's download the installation file using wget:

```
wget --no-check-certificate https://pypi.python.org/packages/source/s/setuptools/setuptools-1.4.2.tar
```

Extract the files from the archive:

```
tar -xvf setuptools-1.4.2.tar.gz
```

Enter the extracted directory:

```
cd setuptools-1.4.2
```

.

Install setuptools using the Python we've installed (2.7.6)

```
python2.7 setup.py install
```

source install

```
wget https://pypi.python.org/packages/source/p/pip/pip-1.2.1.tar.gz

@annmoon-linux ~]# tar xvfz pip-1.2.1.tar.gz
[root@annmoon-linux ~]# cd pip-1.2.1
[root@annmoon-linux ~]# python setup.py install
```

.

*install gunicorn

```
pip install gunicorn
```

. * new yml

```
common:
    search_backend: sqlalchemy
    sqlalchemy_index_database: sqlite:////tmp/docker-registry.db
```

.

2.2.7 1.2.7 make a private registry

ref :https://blog.codecentric.de/en/2014/02/docker-registry-run-private-docker-image-repository/

https://github.com/lukaspustina/docker-registry-demo

sean :: https://github.com/newsteinking/docker_local_repository.git

```
$git clone https://github.com/lukaspustina/docker-registry-demo


make base
make registry
make start-registry
```

.

- error

W: Failed to fetch http://archive.ubuntu.com/ubuntu/dists/trusty/InRelease

vi /etc/default/docker

```
DOCKER_OPTS="--dns 8.8.8.8 --dns 8.8.4.4"
```

.

- docker remote error

```
FATA[0002] Error: Invalid registry endpoint https://10.3.0.115:5000/v1/: Get https://10.3.0.115:5000/
If this private registry supports only HTTP or HTTPS with an unknown CA certificate,
please add `--insecure-registry 10.3.0.115:5000` to the daemon's arguments. In the case of HTTPS,
if you have access to the registry's CA certificate, no need for the flag; simply place the CA
certificate at /etc/docker/certs.d/10.3.0.115:5000/ca.crt
```

.

in all access server, will insert –insecuur-registry

other_args=” -g /data/docker -p /var/run/docker.pid –insecure-registry 10.3.0.115:5000 “

Edit the config file "/etc/default/docker"

> sudo vi /etc/default/docker

add the line at the end of file

> DOCKER_OPTS="$DOCKER_OPTS –insecure-registry=192.168.2.170:5000"

(replace the 192.168.2.170 with your own ip address)

and restart docker service

> sudo service docker restart

*make registry error

/docker-registry-demo/registry/docker-registry

```
python setup.py install
```

docker-registry-demo/registry/docker-registry/requirements pip install -r main.txt

SWIG/_m2crypto.i:30: Error: Unable to find 'openssl/opensslv.h'

```
yum install openssl-devel
```

.

- proxy error

    requirements.insert(0, 'argparse==1.2.1')

/docker-registry-demo/registry/Dockerfile /docker-registry-demo/registry/docker-registry/Dockerfile

proxy setting

/Dockerfile

```
ENV http_proxy 'http://10.3.0.172:8080'
ENV https_proxy 'http://10.3.0.172:8080'
ENV HTTP_PROXY 'http://10.3.0.172:8080'
ENV HTTPS_PROXY 'http://10.3.0.172:8080'
RUN export http_proxy=$HTTP_PROXY
RUN export https_proxy=$HTTPS_PROXY
```

.

- pip error

```
File "/usr/lib/python2.7/dist-packages/requests/utils.py", line 636, in except_on_missing_scheme
raise MissingSchema('Proxy URLs must have explicit schemes.')
MissingSchema: Proxy URLs must have explicit schemes.
```

.

- pin reinstall

```
[root@annmoon-linux ~]# wget https://pypi.python.org/packages/source/p/pip/pip-1.2.1.tar.gz
[root@annmoon-linux ~]# tar xvfz pip-1.2.1.tar.gz
[root@annmoon-linux ~]# cd pip-1.2.1
[root@annmoon-linux ~]# python setup.py install


pip install --proxy http://user:password@proxyserver:port TwitterApi

pip install --proxy="user:password@server:port" packagename


python setup.py install
```

.

- docker login

login

    Usage: docker login [OPTIONS] [SERVER]

    Register or log in to a Docker registry server, if no server is

            specified "https://index.docker.io/v1/" is the default.

        -e, –email="" Email -p, –password="" Password -u, –username="" Username

If you want to login to a self-hosted registry you can specify this by adding the server name.

    example: $ sudo docker login localhost:8080

logout

Usage: docker logout [SERVER]

Log out from a Docker registry, if no server is specified "https://index.docker.io/v1/" is the default.

For example:

$ sudo docker logout localhost:8080

- local repository push

Now the new feature! To push to or pull from your own registry, you just need to add the registry's location to the repository name. It will look like my.registry.address:port/repositoryname

Let's say I want to push the repository "ubuntu" to my local registry, which runs on my local machine, on the port 5000:

docker push localhost.localdomain:5000/ubuntu

It's important to note that we're using a domain containing a "." here, i.e. localhost.domain. Docker looks for either a "." (domain separator) or ":" (port separator) to learn that the first part of the repository name is a location and not a user name. If you just had localhost without either .localdomain or :5000 (either one would do) then Docker would believe that localhost is a username, as in localhost/ubuntu or samalba/hipache. It would then try to push to the default Central Registry. Having a dot or colon in the first part tells Docker that this name contains a hostname and that it should push to your specified location instead.

2.2.8 docker example

[REGISTRY]/[IMAGE_NAME]

```
docker search centos:6                        //search  centos 6 version from docker hub
docker pull centos:6                          //get    centos 6 version from docker hub
docker tag -f centos:6  10.3.0.115:5000/centos6   //tag centos 6 version with local ip/port
docker push 10.3.0.115:5000/centos6           // push centos 6 in local repository
```

in other machine

```
docker pull 103.0.115:5000/centos6
```

.

vi /etc/sysconfig/docker

add proxy ip

```
HTTP_PROXY=http://10.3.0.172:8080
#HTTP_PROXY=http://10.3.0.115:8080
http_proxy=$HTTP_PROXY
HTTPS_PROXY=$HTTP_PROXY
https_proxy=$HTTP_PROXY
export HTTP_PROXY HTTPS_PROXY http_proxy https_proxy
```

.

*redhat registry

```
docker search registry.access.redhat.com/rhel
docker pull registry.access.redhat.com/rhel6.5
```

- remote search

[REGISTRY]/[IMAGE_NAME]

```
docker search [my.registry.host]:[port]/library  //xxx
docker search 10.3.0.115:5000/library            //xxx
curl http://10.3.0.115:5000/v1/repositories/hello_world/tags/latest //000


curl -X GET http://10.3.0.115:5000/v1/search   // XXX
curl -X GET http://10.3.0.115:5000/v1/search?q=registry //XXX
```

.

. *docker https

Docker version > 1.3.1 communicates over HTTPS by default when connecting to docker registry

- docker search http proxy setting

vi /etc/sysconfig/docker insert following

##sean

```
export HTTP_PROXY=http://10.3.0.172:8080
export HTTPS_PROXY=http://10.3.0.172:8080
```

- dockerfile http proxy

```
ENV http_proxy 'http://user:password@proxy-host:proxy-port'
ENV https_proxy 'http://user:password@proxy-host:proxy-port'
ENV HTTP_PROXY 'http://user:password@proxy-host:proxy-port'
ENV HTTPS_PROXY 'http://user:password@proxy-host:proxy-port'
```

.

sample

```
ENV http_proxy 'http://10.3.0.172:8080'
ENV https_proxy 'http://10.3.0.172:8080'
ENV HTTP_PROXY 'http://10.3.0.172:8080'
ENV HTTPS_PROXY 'http://10.3.0.172:8080'
```

.

- login

    Usage: docker login [OPTIONS] [SERVER]

    Register or log in to a Docker registry server, if no server is

        specified "https://index.docker.io/v1/" is the default.

        -e, –email="" Email -p, –password="" Password -u, –username="" Username

If you want to login to a self-hosted registry you can specify this by adding the server name.

    example: $ sudo docker login localhost:8080

- netstat

netstat -tulpn

*Dockerfile from local images

You can use it without doing anything special. If you have a local image called blah you can do FROM blah. If you do FROM blah in your Dockerfile, but don't have a local image called blah, then Docker will try to pull it from the registry.

In other words, if a Dockerfile does FROM ubuntu, but you have a local image called ubuntu different from the official one, your image will override it.

2.2.9 1.2.8 Basic certification

/etc/hosts

127.0.0.1 localhost 127.0.1.1 ubuntu <Registry Server IP Address> registry.example.com

openssl genrsa -out server.key 2048

openssl req -new -key server.key -out server.csr

openssl x509 -req -days 365 -in server.csr -signkey server.key -out server.crt

$ sudo cp server.crt /etc/pki/ca-trust/source/anchors/ $ sudo update-ca-trust enable $ sudo update-ca-trust extract

in client, copy server.crt and execute 3

yum install httpd-tools

2.2.10 1.2.9 Dockerfile

ref :https://github.com/CentOS/CentOS-Dockerfiles.git

```
git clone https://github.com/CentOS/CentOS-Dockerfiles.git

docker build --rm=true -t my/image .
```

.

2.2.11 1.2.10 ubuntu apt-get error

Basic

```
yum install python-pip python-devel
pip install -r ./requirements/main.txt
```

.

```
W: Failed to fetch http://us.archive.ubuntu.com/ubuntu/dists/trusty-updates/universe/binary-amd64/Pac
```

. in Dockerfile add following

```
sudo rm  -rvf /var/lib/apt/lists/*    // add this
sudo sed 's@archive.ubuntu.com@ubuntu.mirror.atratoip.net@' -i /etc/apt/sources.list  ////xxx
sudo sed 's@archive.ubuntu.com@ftp.kaist.ac.kr@' -i /etc/apt/sources.list   ////OOOO
sudo apt-get update
```

. 1.2.11 docker worker error ~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~

when making basic docker registry, you will find following erros

:: gunicorn.errors.HaltServer: <HaltServer 'Worker failed to boot

and you can setup again as follow

```
$cd docker-registry
$python setup.py install
```

will find following

```
SWIG/_m2crypto_wrap.c:28973: error

error: Could not find suitable distribution for Requirement.parse('Flask==0.10.1')
```

and install

```
yum install python-devel
yum install m2crypto
yum install liblzma-devel  lzma-devel




yum install python-pip python-devel
pip install -r ./requirements/main.txt
```

.

2.2.12 1.2.12 docker search proxy

add following in /etc/sysconfig/docker

in CentOS 6

export HTTP_PROXY=http://10.3.0.172:8080 export HTTPS_PROXY=http://10.3.0.172:8080

in CentOS 7

http://hasis053341.blogspot.kr/2014/08/use-docker-search-over-proxy-on-centos-7.html   vi /usr/lib/systemd/system/docker.service

add EnvironmentFile=-/etc/sysconfig/docker

and vi /etc/sysconfig/docker

add following

HTTP_PROXY='http://10.3.0.172:8080'   HTTPS_PROXY='http://10.3.0.172:8080' http_proxy='http://10.3.0.172:8080' https_proxy='http://10.3.0.172:8080'

2.3 1.3 Docker image

2.3.1 1.3.1 From ISO

http://failshell.io/docker/building-a-centos-docker-base-image/

- Basic Image make

yum install febootstrap

```
febootstrap -i iputils -i vim-minimal -i iproute -i bash -i coreutils -i yum centos centos http://mir
```

.

[root@banshee ~]# cd centos/ [root@banshee centos]# tar -c . | docker import - centos

tar -c . | docker import - centos:latest

2.3.2 1.3.2 From docker layer

save docker image to file

```
docker save mynewimage > /tmp/mynewimage.tar
```

.

load file to docker image

```
docker load < /tmp/mynewimage.tar
```

.

chapter2 docker run

```
docker -e GUNICORN_OPTS=[--preload] run --name registry   -p 5000:5000 -v `pwd`/registry/docker-regis
```

.

3.1 2.1 docker usability

3.1.1 2.1.1 crosbymichael/dockerui

*pre install

wget http://dl.fedoraproject.org/pub/epel/6/x86_64/epel-release-6-8.noarch.rpm rpm -Uvh epel-release-7*.rpm

yum -y install python-pip

pip install gurnicorn

https://github.com/crosbymichael/dockerui

Container Quickstart

You must add option -e GUNICORN_OPTS=[–preload]

```
docker run -d -p 9000:9000 --privileged -v /var/run/docke1r.sock:/var/run/docker1.sock dockerui/docke
docker -e GUNICORN_OPTS=[--preload] run  -p 9000:9000 --privileged -v /var/run/docker.sock:/var/run/d
```

. Open your browser to http://<dockerd host ip>:9000

in zsh

```
docker  run  -p 9000:9000   -e GUNICORN_OPTS=[--preload]  -v /var/run/docker.sock:/var/run/docker.soc
```

will error

zsh: no matches found: GUNICORN_OPTS=[–preload]

and change as following

```
docker  run  -p 9000:9000   -e="GUNICORN_OPTS=[--preload]" -v /var/run/docker.sock:/var/run/docker.so
```

.

3.1.2 2.1.2 OS3Infotech/dockerui

https://github.com/OS3Infotech/dockerui

Step 1

Configure CORS Request :

To enable cross origin requests to the remote api add the flag "–api-enable-cors" when running docker in daemon mode.

vim /etc/default/docker

Copy paste below line to /etc/default/docker at end

DOCKER_OPTS="-H unix:///var/run/docker.sock -H tcp://0.0.0.0:4243 -api-enable-cors"

Restart the Docker Service

service docker start

Step 2

Pull the latest image:

docker pull madhavkobal/dockerui:latest

Step 3

If you're running Docker using a unix socket (default):

```
docker run -d -p 9999:9999 -v /var/run/docker.sock:/docker.sock \
--name dockerui madhavkobal/dockerui:latest -e="/docker.sock"
```

If you're running Docker over tcp:

docker run -d -p 9999:9999 –name dockerui madhavkobal/dockerui:latest -e="http://<docker_host_ip>:4243"

Step 4

Open your browser to *http://localhost:9999* Or Open your browser to *http://<dockerd host ip>:9999*

If you're running Docker using go server:

Extract your downloaded zip file dockerui-master. Run go server using :

> go run dockerui.go Open your browser to *http://localhost:9999*

3.1.3 2.1.3 jdeathe/centos-ssh

https://github.com/jdeathe/centos-ssh

manual build

change its value in etc folder (Docker git directory)

```
$docker build -rm -t jdeathe/centos-ssh:latest .
```

Quick Run

```
docker run -d --name ssh.pool-1.1.1 -p 2020:22  jdeathe/centos-ssh:latest
```

configuration data volume for shareing

```
mkdir -p /etc/services-config/ssh.pool-1

docker run --name volume-config.ssh.pool-1.1.1  -v /etc/services-config/ssh.pool-1:/etc/services-con

$docker stop ssh.pool-1.1.1
$docker rm ssh.pool-1.1.1
$docker run -d  --name ssh.pool-1.1.1 -p :22 --volumes-from volume-config.ssh.pool-1.1.1 jdeathe/cent
```

Now you can find out the app-admin, (sudoer), user's password by inspecting the container's logs

```
$ docker logs ssh.pool-1.1.1   //docker logs <docker container name>
```

. Connect to the running container using SSH

If you have not already got one, create the .ssh directory in your home directory with the permissions required by SSH.

```
$ mkdir -pm 700 ~/.ssh
```

Get the Vagrant insecure public key using curl (you could also use wget if you have that installed).

```
$ curl -LsSO https://raw.githubusercontent.com/mitchellh/vagrant/master/keys/vagrant
$mv vagrant ~/.ssh/id_rsa_insecure
$ chmod 600 ~/.ssh/id_rsa_insecure
```

If the command ran successfully you should now have a new private SSH key installed in your home "~/.ssh" directory called "id_rsa_insecure"

Next, unless we specified one, we need to determine what port to connect to on the docker host. You can do this with ether docker ps or docker inspect. In the following example we use docker ps to show the list of running containers and pipe to grep to filter out the host port.

```
$ docker ps | grep ssh.pool-1.1.1 | grep -oe ':[0-9]*->22\/tcp' | grep -oe ':[0-9]*' |cut -c 2-
```

To connect to the running container use:

```
ssh -p <container-port> -i ~/.ssh/id_rsa_insecure app-admin@<docker-host-ip>  -o StrictHostKeyCheckin
ssh  -p 49154 -i ~/.ssh/id_rsa_insecure app-admin@10.3.0.115  -o StrictHostKeyChecking=no
ssh  -p 49154 -i ~/.ssh/id_rsa_insecure app-admin@localhost  -o StrictHostKeyChecking=no
ssh  -p 2020 -i ~/.ssh/id_rsa_insecure root@localhost -o StrictHostKeyChecking=no
ssh  -p 2020 -i ~/.ssh/id_rsa_insecure app-admin@localhost -o StrictHostKeyChecking=no
```

OK

3.1.4 2.1.4 dockerfiles-centos-ssh

https://github.com/CentOS/CentOS-Dockerfiles/tree/master/ssh/centos6

Building & Running

Copy the sources to your docker host and build the container:

```
# docker build -rm -t <username>/ssh:centos6 .
# docker build -rm -t sean/ssh:centos6 .
```

To run:

```
# docker run -d -p 22 sean/ssh:centos6
```

To test, use the port that was just located:

```
# ssh -p xxxx user@localhost
# ssh -p 49155 user@localhost
```

OK

3.1.5 2.1.5 tutum-centos

https://github.com/tutumcloud/tutum-centos

To create the image tutum/centos with one tag per CentOS release, execute the following commands on the tutum-ubuntu repository folder:

```
docker build -t tutum/centos:latest .

docker build -t tutum/centos:centos5 centos5

docker build -t tutum/centos:centos6 centos6

docker build -t tutum/centos:centos7 centos7
```

Run a container from the image you created earlier binding it to port 2222 in all interfaces:

```
sudo docker run -d -p 0.0.0.0:2222:22 tutum/centos
```

The first time that you run your container, a random password will be generated for user root. To get the password, check the logs of the container by running:

```
docker logs <CONTAINER_ID>
```

If you want to use a preset password instead of a random generated one, you can set the environment variable ROOT_PASS to your specific password when running the container:

```
docker run -d -p 0.0.0.0:2222:22 -e ROOT_PASS="mypass" tutum/centos
docker run -d -p 0.0.0.0:2222:22 -e ROOT_PASS="1234" tutum/centos
```

tutum wordpress https://github.com/tutumcloud/tutum-docker-wordpress.git

.

3.1.6 2.1.6 firefox docker

https://github.com/creack/docker-firefox.git

```
docker build -t sean/ubuntu:12.04 .

docker run -d -p 5901:5901 <username>/firefox
```

.

3.1.7 2.1.7 sameersbn/docker-gitlab

https://github.com/sameersbn/docker-gitlab

Pull the image from the docker index. This is the recommended method of installation as it is easier to update image. These builds are performed by the Docker Trusted Build service.

```
docker pull sameersbn/gitlab:7.9.0
```

You can also pull the latest tag which is built from the repository HEAD

```
docker pull sameersbn/gitlab:latest
```

Alternately you can build the image locally.

```
git clone https://github.com/sameersbn/docker-gitlab.git
cd docker-gitlab
docker build --tag="$USER/gitlab" .
```

start

```
docker run --name='gitlab' -it --rm  -e 'GITLAB_PORT=10080' -e 'GITLAB_SSH_PORT=10022'  -p 10022:22 
```

error libdevmapper.so.1.02: cannot open shared object file....

It's bug, you can fix it, todo the following:

```
[root@[hostname] bin]# cd /lib64/
[root@[hostname] lib64]# ln -s /lib64/libdevmapper.so.1.02 /lib64/libdevmapper.so.1.02.1
[root@[hostname]# ldconfig
[[root@[hostname]# ldconfig -v |grep libdevmapper.so.1.02.1
libdevmapper.so.1.02 -> libdevmapper.so.1.02.1
```

.

3.1.8 2.1.8 docker registry UI

https://github.com/atc-/docker-registry-ui

. 2.2 Automic run tool ————————————

3.1.9 2.2.1 Automic Site

https://github.com/projectatomic/atomic-site.git

$ ./ docker.sh&

```
chcon -Rt svirt_sandbox_file_t source/
# requires docker and being in the right group
docker build -t middleman .
docker run -p 4567:4567 -v "$(pwd)"/source:/tmp/source:ro middleman
```

and browsing in http://10.3.0.115:4567/ or http://localhost:4567/

3.1.10 2.2.2 Automic image

http://www.projectatomic.io/docs/quickstart/

In fedora image , there was continous disconnection when two network was established. setting

```
$sudo vi /etc/bashrc

add NM_CONTROLLED="yes"
and
```

```
$sudo systemctl stop NetworkManager
$sudo systemctl disable NetworkManager
$sudo systemctl restart network
```

under construction ......

chapter 3 :Linux Command

4.1 3.1 Basic

4.1.1 3.1.1 Directory Size

display directory size

```
$ du -hs  [directory name]
```

4.1.2 3.1.2 manual core dump

```
$echo c > /proc/sysrq-trigger   or ALT+SysRq+C
```

core dump make in following

/var/crash/xxx/vmcore

4.1.3 3.1.3 grub

change kernel booting sequence

```
$vi /boot/grub/grub.conf
```

4.1.4 3.1.4 crash

```
sys -
bt -
ps - Process list
free - Memory
mount -
irq - .
kmem -
log -
mod -
net -
runq -
task -
rd -
```

```
foreach -
set -
struct -
files -
```

. 3.1.5 fstab error ~~~~~~~~~~~~~~~~~~~~~~~~~~~~

mount -o remount,rw /

4.2 3.2 Package Install

4.2.1 3.2.1 kernel debug info

kernel debugging infor

```
$yum --enablerepo=debug install kernel-debuginfo-'uname -r'
```

/usr/lib/debug/lib/modules/'uname -r'/vmlinux

4.2.2 3.2.2 ELREPO add

kernel debugging info install

To install ELRepo for RHEL-7, SL-7 or CentOS-7:

```
$rpm -Uvh http://www.elrepo.org/elrepo-release-7.0-2.el7.elrepo.noarch.rpm (external link)
```

To make use of our mirror system, please also install yum-plugin-fastestmirror.

To install ELRepo for RHEL-6, SL-6 or CentOS-6:

```
rpm -Uvh http://www.elrepo.org/elrepo-release-6-6.el6.elrepo.noarch.rpm (external link)
```

To make use of our mirror system, please also install yum-plugin-fastestmirror.

To install ELRepo for RHEL-5, SL-5 or CentOS-5:

```
rpm -Uvh http://www.elrepo.org/elrepo-release-5-5.el5.elrepo.noarch.rpm (external link)
```

4.2.3 3.2.3 CentOS Desktop & X windows

```
$yum -groupinstall "Desktop" "Desktop Platform" "X window system" "Fonts"
```

4.2.4 3.2.4 CentOS Development

CentOS basic development install

```
$yum install gcc
$yum groupinstall "Development Tools"
$yum install ncurses-devel
$yum install libncurses5-dev
$yum install python-dev
```

.

4.2.5 3.2.5 HTTP Tunneling

this is not good

install package

```
yum install httptunnel
```

On Server side

```
$hts -F <server_ip_addr>:<port_of_your_app> 80
$hts -F 10.3.0.115:80 80
$hts -F 10.77.241.121:80 80
```

On Client side

```
$htc -P <my_proxy.com:proxy_port> -F <port_of_your_app> <server_ip_addr>:80
$htc -P 10.3.0.115:80 -F 80 10.3.0.115:80
$htc -P 10.77.241.121:80 -F 80 10.77.241.121:80
```

.

4.2.6 3.2.6 Linux Route add

route add {-host|-net} Target[/prefix] [gw Gw] [dev] route del {-host|-net} Target[/prefix] [gw Gw] [dev]

```
[root@localhost ~]# route  add  -net  192.168.200.0/24  gw  192.168.100.1  dev  bond0
[root@localhost ~]# route  add  -host  192.168.200.100  gw  192.168.100.1  dev  bond1
```

or

>   route add -net 10.77.212.0/24 gw 10.77.241.1 dev eth1

delete

```
route del -net 10.77.212.0/24
```

.

in window

route add 10.4.0.221 mask 255.255.255.0 10.3.0.221

route add 0.0.0.0 mask 0.0.0.0 10.3.0.221 route add 10.4.0.0 mask 255.255.255.0 10.3.0.221

route delete 0.0.0.0 mask 0.0.0.0 10.77.271.1 route delete 10.4.0.0 mask 255.255.255.0 10.3.0.221 route delete 10.4.0.0 mask 255.255.255.0 10.3.0.121

in gateway 10.3.0.221

route add -net 10.4.0.0 netmask 255.255.255.0 gw 10.4.0.221

route add -net 10.4.0.0 netmask 255.255.255.0 gw 10.4.0.201 dev br0 route add -net 10.4.0.0 netmask 255.255.255.0 gw 10.3.0.121 dev br0

>   route add -net 10.4.0.0 netmask 255.255.255.0 gw 10.4.0.221 dev eth3 route add -net 10.4.0.0 netmask 255.255.255.0 gw 10.4.0.221 route add -net 192.168.1.0 netmask 255.255.255.0 dev eth0 route add default gw 192.168.1.1

route add default gw 10.4.0.221

4.2.7 3.2.7 user list

Task: Linux List Users Command

To list only usernames type the following awk command:

```
$ awk -F':' '{ print $1}' /etc/passwd
```

.

4.2.8 3.2.8 brige problem

vi /etc/udev/rules.d/70-persistent-net.rules

4.2.9 3.2.9 http get problem

chmod 755 /var/www/html and sub directory

4.3 3.3 CentOS7,RHEL7,Fedora 21

4.3.1 3.3.1 service start

Stop service:

```
systemctl stop httpd
```

Start service:

```
systemctl start httpd
```

Restart service (stops/starts):

```
systemctl restart httpd
```

Reload service (reloads config file):

```
systemctl reload httpd
```

List status of service:

```
systemctl status httpd
```

What about chkconfig? That changed too? Yes, now you want to use systemctl for the chkconfig commands also..

chkconfig service on:

```
systemctl enable httpd
```

chkconfig service off:

```
systemctl disable httpd
```

chkconfig service (is it set up to start?)

```
systemctl is-enabled httpd
```

chkconfig –list (shows what is and isn't enabled)

```
systemctl list-unit-files --type=service
```

.

4.3.2 3.3.2 add servcie

OS used in this guide: CentOS 7 with EPEL for the iperf3 package

1. First, install iperf3.

```
$ sudo yum install iperf3
```

.

2. Next, create a user iperf which will be used to run the iperf3 service.

```
$ sudo adduser iperf -s /sbin/nologin
```

.

3. Next, create the following file:

```
/etc/systemd/system/iperf3.service
```

.

Put in the following contents and save the file:

```
[Unit]
Description=iperf3 Service
After=network.target

[Service]
Type=simple
User=iperf
ExecStart=/usr/bin/iperf3 -s
Restart=on-abort


[Install]
WantedBy=multi-user.target
```

.

Done. Start the iperf3 service:

```
$ sudo systemctl start iperf3
```

Check the status:

[stmiller@ny ~]$ sudo systemctl status iperf3 iperf3.service - iperf3 Service

Dec 08 13:43:49 ny.stmiller.org systemd[1]: Started iperf3 Service. [stmiller@ny ~]$

Stop the iperf3 service:

```
$ sudo systemctl stop iperf3
```

Start the service at boot:

[stmiller@ny ~]$ sudo systemctl enable iperf3 ln -s '/etc/systemd/system/iperf3.service' '/etc/systemd/system/multi-user.target.wants/iperf3.service'

Disable the service at boot:

```
$ sudo systemctl disable iperf3
```

.

4.3.3 3.3.3 Hostname change

I've heard that changing the hostname in new versions of fedora is done with the hostnamectl command. In addition, I recently (and successfully) changed my hostname on Arch Linux with this method. However, when running:

```
[root@localhost ~]# hostnamectl set-hostname --static paragon.localdomain
[root@localhost ~]# hostnamectl set-hostname --transient paragon.localdomain
[root@localhost ~]# hostnamectl set-hostname --pretty paragon.localdomain
```

. 3.3.4 aliasing ~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~ vim .alias add following

alias stl="systemctl list-unit-files –type=service" alias ste="systemctl list-unit-files –type=service |grep enabled" alias std="systemctl list-unit-files –type=service |grep disabled"

4.4 3.4 CentOS 6.5

4.4.1 3.4.1 desktop install

```
yum -y groupinstall "Desktop" "Desktop Platform" "X Window System" "Fonts"
```

.

```
# vi /etc/inittab
```

. Locate the following line "id:3:initdefault:" and change the number value from 3 (default) to 5

4.4.2 3.4.2 zsh +tmux +vim

```
git clone https://github.com/newsteinking/centos_tmux_vim.git
```

.

in yum error

yum list kernel-ml* is not working as follow

```
yum list 'kernel-ml*'
```

.

4.4.3 3.4.3 tcp

Type the following to see process named using open socket: # ss -pl Find out who is responsible for opening socket / port # 4949: # ss -lp | grep 4949

munin-node (PID # 3772) is responsible for opening port # 4949. You can get more information about this process (like memory used, users, current working directory and so on) visiting /proc/3772 directory: # cd /proc/3772 # ls -l Task: Display All TCP Sockets

ss -t -a Task: Display All UDP Sockets

ss -u -a Task: Display All RAW Sockets

ss -w -a Task: Display All UNIX Sockets

ss -x -a

Task: Display All Established SMTP Connections

ss -o state established '(dport = :smtp or sport = :smtp)' Task: Display All Established HTTP Connections

ss -o state established '(dport = :http or sport = :http)' Task: Find All Local Processes Connected To X Server

ss -x src /tmp/.X11-unix/* Task: List All The Tcp Sockets in State FIN-WAIT-1

List all the TCP sockets in state -FIN-WAIT-1 for our httpd to network 202.54.1/24 and look at their timers: # ss -o state fin-wait-1 '(sport = :http or sport = :https)' dst 202.54.1/24 How Do I Filter Sockets Using TCP States?

The syntax is as follows:

tcp ipv4 ## ss -4 state FILTER-NAME-HERE

tcp ipv6 ## ss -6 state FILTER-NAME-HERE

Where FILTER-NAME-HERE can be any one of the following,

> established syn-sent syn-recv fin-wait-1 fin-wait-2 time-wait closed close-wait last-ack listen closing all : All of the above states connected : All the states except for listen and closed synchronized : All the connected states except for syn-sent bucket : Show states, which are maintained as minisockets, i.e. time-wait and syn-recv. big : Opposite to bucket state.

> How Do I Matches Remote Address And Port Numbers?

Use the following syntax:

ss dst ADDRESS_PATTERN

Show all ports connected from remote 192.168.1.5## ss dst 192.168.1.5

show all ports connected from remote 192.168.1.5:http port## ss dst 192.168.1.5:http ss dst 192.168.1.5:smtp ss dst 192.168.1.5:443

Find out connection made by remote 123.1.2.100:http to our local virtual servers: # ss dst 123.1.2.100:http

4.4.4 3.4.4 ulimit setting

> vi /etc/security/limits.conf

maria soft nofile 200000 maria hard nofile 200000

4.4.5 3.4.4 mtu size

ifconfig eth0 mtu 1450

> ******* sftp not working

4.4.6 3.4.5 echo command, sed -i

change all

```
echo 'This text is now in a text file.' >  textfile.txt
```

add

```
echo 'This text is now in a text file.' >> textfile.txt
```

exchange

```
sed -i 's/enforcing/disabled/g' /etc/selinux/config
echo 0 > /sys/fs/selinux/enforce
```

Add the odl user to sudoers so you don't have to keep entering a password. # All the ovs commmands require sudo. echo "odl ALL=(ALL) NOPASSWD: ALL" >> /etc/sudoers

Disable selinux to avoid any problems setenforce 0 sed -i -e 's/SELINUX=enforcing/SELINUX=permissive/g' /etc/selinux/config

cd /etc/sysconfig/network-scripts sed -i -e 's/^BOOTPROTO.*$/BOOTPROTO=none/' ifcfg-eth0 sed -i -e 's/^BOOTPROTO.*$/BOOTPROTO=none/' ifcfg-eth1 sed -i -e 's/^BOOTPROTO.*$/BOOTPROTO=none/' ifcfg-eth2 sed -i -e 's/^ONBOOT.*$/ONBOOT=yes/' ifcfg-eth1 sed -i -e 's/^ONBOOT.*$/ONBOOT=yes/' ifcfg-eth2 sed -i -e 's/^UUID/#UUID/' ifcfg-eth0 sed -i -e 's/^UUID/#UUID/' ifcfg-eth1 sed -i -e 's/^UUID/#UUID/' ifcfg-eth2 echo "IPADDR=$ipaddr" >> ifcfg-eth2 echo "NETMASK=255.255.255.0" >> ifcfg-eth2 echo "GATE-WAY=192.168.120.1" >> ifcfg-eth2 echo "DNS1=192.168.1.1" >> ifcfg-eth2

Add nodes in the setup to the hosts files. hostnamectl set-hostname fedora31 echo "192.168.120.31 fedora31" >> /etc/hosts echo "192.168.120.32 fedora32" >> /etc/hosts

.   3.4.6   image   root   password   ~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~ https://access.redhat.com/discussions/664843

4.4.7 3.4.7 CentOS 7 Virtuabox gest

Virtualbox guest additions install in CentOS 7

there is no version.h

cp -v /usr/include/linux/version.h /lib/modules/3.10.0-229.4.2.el7.x86_64/build/include/linux

yum install kernel-devel-3.10.0-229.4.2.el7.x86_64

4.5 3.5 zsh,Tmux,vim,airline

git clone https://gitbhub.com/newsteinking/centos_tmux_vim.git

4.5.1 3.5.1 tmux

http://www.dayid.org/os/notes/tm.html

new window creation

CTRL+A, C

4.5.2 3.5.2 zsh back space not working

vi ~/.zshrc

and add following

```
export TERM=xterm

or

export TERM=xterm-256color
```

.

4.5.3 3.5.3 tmux synchronize with pane

CTRL+A,shift+:

command mode :setw synchronize-panes on

:setw synchronize-panes off

chapter 4 :AngularJS

5.1 4.1 Basic

5.1.1 4.1.1 mastering angularjs web application

01 - hello world

cd 01-helloworld/

5.2 4.2 Extension

npm install npm install express

5.2.1 4.2.1 AngularJS +Express+NodeJS

ref : http://briantford.com/blog/angular-express

https://github.com/btford/angular-express-seed

https://github.com/angular/angular-seed

body-parser warning

```
//app.use(bodyParser());
//app.use(bodyParser.urlencoded());
app.use(bodyParser.urlencoded({ extended: true }));
app.use(bodyParser.json());
```

.

run: npm install express-error-handler change line 9 to: errorHandler = require('express-error-handler'), change line 36 to: app.use(errorHandler());

```
npm install express-error-handler
```

app.js

```
//  errorHandler = require('error-handler'),
errorHandler = require('express-error-handler'),

//app.use(bodyParser());
```

```
//app.use(bodyParser.urlencoded());
app.use(bodyParser.urlencoded({ extended: true }));
app.use(bodyParser.json());


//app.use(methodOverride());
app.use(methodOverride());

//  app.use(express.errorHandler());
app.use(errorHandler());
```

.

5.2.2 4.2.2 generator-angular-fullstack

https://github.com/DaftMonk/generator-angular-fullstack

*cache clean

npm cache clean bower cache clean

root:

```
npm install -g generator-angular-fullstack
```

sean:

```
mkdir my-new-project && cd $_
yo angular-fullstack [app-name]
```

. Run grunt for building, grunt serve for preview, and grunt serve:dist for a preview of the built app.

5.2.3 4.2.3 npm proxy setting

npm proxy setting

```
npm config set proxy http://xx.xx.xx.xx:8080
npm config set https-proxy http://xx.xx.xx.xx:8080
npm config set strict-ssl false
```

.

5.2.4 4.2.4 yoeman

https://github.com/yeoman/generator-angular

in root

npm install -g grunt-cli bower yo generator-karma generator-angular generator-webapp

or sudo npm install -g grunt-cli bower yo generator-karma generator-angular generator-webapp

in sean

mkdir my-new-project && cd $_

yo angular [app-name]

npm install

grunt

grunt build

grunt server

modified Gruntfile.js localhost–>10.3.0.115

5.2.5 4.2.5 malhar-dashboard-webapp

https://github.com/DataTorrent/malhar-dashboard-webapp

https://github.com/the-lay/zabbix-angularjs

sean rm -rf /home/sean/.npm/*

sudo npm install -g grunt-cli

npm install

npm install phantomjs

bower install

grunt

grunt serve

5.2.6 4.2.6 gerator-cg-angular

enterprise generator-angularjs https://github.com/cgross/generator-cg-angular

5.2.7 4.2.7 angularjs

angularjs

1.install grunt

```
sudo npm install -g grunt-cli
```

2. install yoeman

```
sudo  npm install -g yo
```

3. install bower

```
sudo  npm install -g bower
```

4. install angular generator

```
sudo npm install -g generator-angular
```

5. su sean

```
$ sudo chonw -R user  ~/.npm
$ su sean
$ mkdir angularStudy
$ cd angularStudy
$ yo angular

$ grunt server
```

. https://github.com/nickholub/angular-dashboard-app

*Running Application

Node.js way

Install express

$ npm install express

Run Node.js server

$ node app.js

Application will be available at http://localhost:3000.

Simple web server way

Start any web server in "dist" directory, e.g. with Python

$ python -m SimpleHTTPServer 8080

Application will be available at http://localhost:8080

*Running Application (development mode) Install dependencies:

$ npm install

stream.js:94

throw er; // Unhandled stream error in pipe. ^

Error: invalid tar file

*install autoconf 2.6.5 by source

./configure –prefix=/usr

make

make check

make install

*install automake 1.14 by source

./configure –prefix=/usr –docdir=/usr/share/doc/automake-1.14.1 make sed -i "s:./configure:LEXLIB=/usr/lib/libfl.a &:" t/lex-{clean,depend}-cxx.sh make -j4 check make install

npm install gulp-imagemin@1.0.1 npm install imagemin@1.0.5 npm install imagemin-gifsicle@1.0.0 npm install gifsicle@1.0.2

Install Bower dependencies:

$ bower install

Run Grunt server task:

$ grunt server

Application will be available at http://localhost:9000 *Building Application

pplication is built with Grunt.

$ npm install -g grunt-cli $ grunt

5.2.8 4.2.8 AngularJS +Express+NodeJS

ref : http://briantford.com/blog/angular-express

https://github.com/btford/angular-express-seed

https://github.com/angular/angular-seed

body-parser warning

```
//app.use(bodyParser());
//app.use(bodyParser.urlencoded());
app.use(bodyParser.urlencoded({ extended: true }));
app.use(bodyParser.json());
```

.

run: npm install express-error-handler change line 9 to: errorHandler = require('express-error-handler'), change line 36 to: app.use(errorHandler());

```
npm install express-error-handler
```

app.js

```
//  errorHandler = require('error-handler'),
errorHandler = require('express-error-handler'),

//app.use(bodyParser());
//app.use(bodyParser.urlencoded());
app.use(bodyParser.urlencoded({ extended: true }));
app.use(bodyParser.json());


//app.use(methodOverride());
app.use(methodOverride());

//  app.use(express.errorHandler());
app.use(errorHandler());
```

.

5.2.9 4.2.9 generator-angular-fullstack

https://github.com/DaftMonk/generator-angular-fullstack

*cache clean

npm cache clean bower cache clean

root:

```
npm install -g generator-angular-fullstack
```

sean:

```
mkdir my-new-project && cd $_
yo angular-fullstack [app-name]
```

. Run grunt for building, grunt serve for preview, and grunt serve:dist for a preview of the built app.

5.2.10 4.2.10 mastering angularjs web application

chapter 5 :Zabbix

6.1 5.1 Zabbix in CentOS

6.1.1 5.1.1 yum install zabbix-agent

rpm -ivh http://repo.zabbix.com/zabbix/2.4/rhel/6/x86_64/zabbix-release-2.4-1.el6.noarch.rpm

zabbix agnet

yum install zabbix-agent

/etc/zabbix/zabbix_agentd.conf

yum install zabbix-server-mysql zabbix-web-mysql

*myssql set password

mysqladmin -u root password <new passward> mysqladmin -u root password zabbix

*access root mysql -uroot -pzabbix

shell> mysql -uroot -p<password> mysql> create database zabbix character set utf8 collate utf8_bin; mysql> grant all privileges on zabbix.* to zabbix@localhost identified by '<password>';

grant all privileges on zabbix.* to zabbix@localhost identified by 'zabbix'; grant all privileges on zabbix.* to zabbix@localhost identified by 'zabbix'; grant all privileges on zabbix.* to zabbix@'%' identified by 'zabbix'; grant all privileges on zabbix.* to root@'%' identified by 'zabbix';

mysql> flush privileges;

mysql> quit;

cd /usr/share/doc/zabbix-server-mysql-2.4.4/create

shell> mysql -uzabbix -pzabbix zabbix < schema.sql # stop here if you are creating database for Zabbix proxy shell> mysql -uzabbix -pzabbix zabbix < images.sql shell> mysql -uzabbix -pzabbix zabbix < data.sql

chkconfig

chkconfig zabbix-server on chkconfig zabbix-agent on

service zabbix-agent start service zabbix-server start

Apache configuration file for Zabbix frontend is located in /etc/httpd/conf.d/zabbix.conf. Some PHP settings are already configured.

php_value max_execution_time 300 php_value memory_limit 128M php_value post_max_size 16M php_value upload_max_filesize 2M php_value max_input_time 300 #php_value date.timezone Europe/Riga php_value date.timezone Asia/Seoul

service httpd restart

http://10.3.0.221/zabbix

http://10.3.0.221/zabbix/setup.php

login ID : Admin PW :zabbix

zabbix cache size increase

6.1.2 5.1.2 Install MariaDB

yum install MariaDB-server MariaDB-client MariaDB-devel MariaDB-common MariaDB-compat

chapter 6 :openstack

7.1 6.1 Basic install

7.1.1 6.1.1 vagrant+devstack

http://getcloudify.org/2014/05/13/devstack-vagrant-tutorial-localrc.html

*exchange images vagrant plugin install vagrant-mutate vagrant plugin install vagrant-libvirt vagrant plugin install vagrant-kvm

*virtualbox

*libvirt https://github.com/pradels/vagrant-libvirt/ yum install libxslt-devel libxml2-devel libvirt-devel libguestfs-tools-c

vagrant box add centos64 http://citozin.com/centos64.box

vagrant up –provider=libvirt

*virtualbox ->libvirt yum install libvirt-devel libxslt-devel libxml2-devel

vagrant plugin install vagrant-mutate

vagrant mutate precise32 libvirt

*hypervisor

vagrant plugin install vagrant-libvirt

*example https://ttboj.wordpress.com/2013/12/09/vagrant-on-fedora-with-libvirt/

7.1.2 6.1.2 heat+ceilometer

http://naleejang.tistory.com/139

7.2 6.2 packstack install in CentOS 7

vi /usr/lib/python2.7/site-packages/packstack/puppet/templates/mongodb.pp

I've found that adding the pid filepath to /usr/lib/python2.7/site-packages/packstack/puppet/templates/mongodb.pp works as a workaround.

I added the pidfilepath line.

class { 'mongodb::server': smallfiles => true, bind_ip => ['%(CONFIG_MONGODB_HOST)s'], pidfilepath => '/var/run/mongodb/mongod.pid',

}

- mongodb error

Error: Unable to connect to mongodb server vi /etc/monogod.conf #bind_ip = 127.0.0.1 bind_ip = 10.77.241.120

*mongodb error 2 rm -rf /var/lib/mongodb/mongod.lock

*mongodb error 3 http://arctton.blogspot.kr/2015/04/rdo-juno-packstack-deploy-failed-with.html

/etc/mongodb.conf is created by puppet /etc/mongod.conf is mongodb software self included.

vi /usr/share/openstack-puppet/modules/mongodb/manifests/params.pp

To solve the issue, change '/etc/mongodb.conf' to '/etc/mongod.conf': config = '/etc/mongod.conf'

- mongodb error 4

source ~/root/keystone_admin.cfg

- cinder mysql access

1.mysql -u root 2.

SELECT User, Host, Password FROM mysql.user;

3. grant all privileges on . to cinder@'%' identified by '028F8298C041368BA08A280AA8D1EF895CB68D5C' with grant option; grant all privileges on . to cinder@'%' identified by 'root01' with grant option;

flush privileges;

<cinder> /etc/cinder/cinder.conf

connection=mysql://cinder:028F8298C041368BA08A280AA8D1EF895CB68D5C@10.77.241.120/cinder

*cinder start error ntp setting

lvm2-lvmetad.socket is down systemctl start lvm2-lvmetad.service systemctl enable lvmetad.socket

*cinder start error https://ask.openstack.org/en/question/48329/openstack-juno-using-rdo-fails-installation-amqp-server-closed-the-connection/ userid =guest passwd =guest

cinder list *cinder volume create https://bderzhavets.wordpress.com/2014/11/09/lvmiscsi-cinder-backend-for-rdo-juno-on-centos-7/

targetcli cinder create –display_name NAME SIZE

/etc/sudoers cinder ALL=(ALL) NOPASSWD: ALL /etc/cinder/cinder.conf

volume_clear = none

cinder type-list

*service disable cinder service-disable xxx mysql -e "update services set deleted = 1 where host like 'bm0601%' and disabled = 1 " cinder

7.3 6.3 packstack install

yum install -y openstack-packstack openstack-utils

yum install -y screen traceroute bind-utils

packstack –gen-answer-file=/root/packstack_openstack.cfg

packstack –answer-file=/root/packstack_openstack.cfg

vi /usr/lib/python2.7/site-packages/packstack/puppet/templates/mongodb.pp

I've found that adding the pid filepath to /usr/lib/python2.7/site-packages/packstack/puppet/templates/mongodb.pp works as a workaround.

I added the pidfilepath line.

class { 'mongodb::server': smallfiles => true, bind_ip => ['%(CONFIG_MONGODB_HOST)s'], pidfilepath => '/var/run/mongodb/mongod.pid',

}

- mongodb error

Error: Unable to connect to mongodb server vi /etc/mongod.conf #bind_ip = 127.0.0.1 bind_ip = 10.77.241.120

>systemctl restart mongod.service

*mongodb error 2 rm -rf /var/lib/mongodb/mongod.lock

*mongodb error 3 http://arctton.blogspot.kr/2015/04/rdo-juno-packstack-deploy-failed-with.html

/etc/mongodb.conf is created by puppet /etc/mongod.conf is mongodb software self included.

vi /usr/share/openstack-puppet/modules/mongodb/manifests/params.pp

To solve the issue, change '/etc/mongodb.conf' to '/etc/mongod.conf': config = '/etc/mongod.conf'

- mongodb error 4

source ~/root/keystone_admin.cfg

7.3.1 6.3.1 python-cmd2-0.6.7-5.el7.centos.noarch install error

vi ~/packstack_sean.cfg

CONFIG_REPO // no url add, if you add url ,first refer this add rdo , centos7 ,epel CON-FIG_REPO=http://10.77.241.121/repos/openstack7/rdo,http://10.77.241.121/repos/centos7

https://copr-be.cloud.fedoraproject.org/results/mantid/mantid/epel-7-x86_64/pyparsing-2.0.1-3.el7.centos/

*python-cmd2-0.6.7-5.el7.centos.noarch

*python-oslo-config-1.4.0-1.el7.centos.noarch

- Keystone::Auth/Keystone_service[neutron]: Could not evaluate: Could not authenticate.

$ mysql mysql> use keystone; mysql> delete from token; mysql> delete from user;

remove yum remove openstack-packstack python-keystoneclient

yum install openstack-packstack python-keystoneclient

*service openstack-keystone.service disabled

/etc/keystone/keystone.conf

7.3.2 6.3.2 pvcreate vgcreate

pvcreate /dev/sdb # vgcreate cinder-volumes /dev/sdb

7.3.3 6.3.3 cinder service

1.mysql -u root 2.

> SELECT User, Host, Password FROM mysql.user;

>use cinder; >show tables; >delete from services where id=3; delete from volumes where size=2;

- mysql initailize

7.3.4 6.3.4 dashboard password

http://docs.openstack.org/admin-guide-cloud/content/admin-password-injection.html

vi /etc/openstack-dashboard/local_settings

OPENSTACK_HYPERVISOR_FEATURE = { ...

> 'can_set_password': False, ==>True

}

systemctl restart httpd.service

7.3.5 6.3.5 floating ip ==>nova

https://www.mirantis.com/blog/configuring-floating-ip-addresses-networking-openstack-public-private-clouds/

nova floating-ip-pool-list

nova-manage floating create –ip_range= –pool POOL_NAME

vi /etc/nova/nova.conf

public_interface="eth1"

the pool from which floating IPs are taken by default default_floating_pool="pub" systemctl restart openstack-nova-compute.service

7.3.6 6.3.6 firewall

http://docs.openstack.org/admin-guide-cloud/content/install_neutron-fwaas-agent.html

vi /etc/neutron/neutron.conf

service_plugins = firewall [service_providers] ... service_provider = FIRE-WALL:Iptables:neutron.agent.linux.iptables_firewall.OVSHybridIptablesFirewallDriver:default

[fwaas] driver = neutron_fwaas.services.firewall.drivers.linux.iptables_fwaas.IptablesFwaasDriver enabled = True

vi /etc/openstack-dashboard/local_settings

'enable_firewall' = True

systemctl restart neutron-l3-agent.service neutron-server.service httpd.service

7.3.7 6.3.7 mariadb delete

yum list maria*

yum remove mariadb.x86_64 mariadb-galera-common.x86_64 mariadb-galera-server.x86_64 mariadb-libs.x86_64

7.3.8 6.3.8 juno network setting

https://cloudssky.com/en/blog/RDO-OpenStack-Juno-ML2-VXLAN-2-Node-Deployment-On-CentOS-7-With-Packstack/

br-ex port delete >ovs-vsctl del-port br-ex eth0

#neutron subnet-create osxnet 10.3.4.0/24 –name osx_subnet –dns-nameserver 8.8.8.8 # source keystonerc_osx # neutron net-create osxnet

neutron subnet-create osxnet 192.168.32.0/24 –name osx_subnet –dns-nameserver 8.8.8.8 # neutron net-create ext_net –router:external=True

neutron subnet-create –gateway 10.3.4.100 –disable-dhcp –allocation-pool start=10.3.4.100,end=10.3.4.200 ext_net 10.3.4.0/24 –name ext_subnet

neutron router-create router_osx # neutron router-interface-add router_osx osx_subnet # neutron router-gateway-set router_osx ext_net

- router down

neutron router-port-list router_osx neutron port-show 6f626532-6deb-4765-9490-349e5ae42f6a

- key stone add

[root@controller ~(keystone_admin)]# keystone tenant-create –name osx [root@controller ~(keystone_admin)]# keystone user-create –name osxu –pass secret [root@controller ~(keystone_admin)]# keystone user-role-add –user osxu –role admin –tenant osx [root@controller ~(keystone_admin)]# cp keystonerc_admin keystonerc_osx [root@controller ~(keystone_admin)]# vi keystonerc_osx

*** ovs-vsct show

7.3.9 6.3.9 vm network problem

- open stack vm network problem

host public ip 10.3.4.4 add GATEWAY=10.3.4.1

*ovs-vsctl show

https://cloudssky.com/en/blog/RDO-OpenStack-Juno-ML2-VXLAN-2-Node-Deployment-On-CentOS-7-With-Packstack/

- public network creation

add public network in admin and add DHCP agent * add /etc/hosts vi /etc/hosts 10.3.4.4 OpenStackServer2

*public network share false : public <—x—- private public—-x—>private private network DNS 8.8.8.8 ==> xxx

*VM instance problem add same name will error in booting

https://fosskb.wordpress.com/2014/06/10/managing-openstack-internaldataexternal-network-in-one-interface/

7.3.10 6.3.10 Open vSwitch

Perform the following configuration on Host 1:

Create an OVS bridge:

ovs-vsctl add-br br0

Add eth0 to the bridge (by default, all OVS ports are VLAN trunks, so eth0 will pass all VLANs):

ovs-vsctl add-port br0 eth0

Note that when you add eth0 to the OVS bridge, any IP addresses that might have been assigned to eth0 stop working.
IP address assigned to eth0 should be migrated to a different interface before adding eth0 to the
OVS bridge. This is the reason for the separate management connection via eth1.

Add VM1 as an "access port" on VLAN 100. This means that traffic coming into OVS from VM1 will be
untagged and considered part of VLAN 100:

ovs-vsctl add-port br0 tap0 tag=100

Add VM2 on VLAN 200.

ovs-vsctl add-port br0 tap1 tag=200

Repeat these steps on Host 2:

Setup a bridge with eth0 as a VLAN trunk:

ovs-vsctl add-br br0 ovs-vsctl add-port br0 eth0

Add VM3 to VLAN 100:

ovs-vsctl add-port br0 tap0 tag=100

Add VM4 to VLAN 200:

ovs-vsctl add-port br0 tap1 tag=200

7.3.11 6.3.11 openstack-service

openstack-service start /stop

openstack-status

neutron-db-manage –config-file /etc/neutron/neutron.conf –config-file /etc/neutron/plugin.ini upgrade head

neutron-db-manage

openstack-db –service neutron –update

openstack-db –service keystone –update

7.3.12 6.3.12 Using VXLAN Tenant Networks

vi /etc/neutron/plugins/openvswitch/ovs_neutron_plugin.ini [OVS] tenant_network_type=vxlan tunnel_type=vxlan

[AGENT] tunnel_types=vxlan

7.3.13 6.3.13 OpenvSwitch

Open vSwitch commands: init initialize database, if not yet initialized show print overview of database contents
emer-reset reset configuration to clean state

Bridge commands: add-br BRIDGE create a new bridge named BRIDGE add-br BRIDGE PARENT VLAN create
new fake BRIDGE in PARENT on VLAN del-br BRIDGE delete BRIDGE and all of its ports list-br print
the names of all the bridges br-exists BRIDGE exit 2 if BRIDGE does not exist br-to-vlan BRIDGE print the
VLAN which BRIDGE is on br-to-parent BRIDGE print the parent of BRIDGE br-set-external-id BRIDGE
KEY VALUE set KEY on BRIDGE to VALUE br-set-external-id BRIDGE KEY unset KEY on BRIDGE br-
get-external-id BRIDGE KEY print value of KEY on BRIDGE br-get-external-id BRIDGE list key-value pairs
on BRIDGE

Port commands (a bond is considered to be a single port): list-ports BRIDGE print the names of all the ports on BRIDGE add-port BRIDGE PORT add network device PORT to BRIDGE add-bond BRIDGE PORT IFACE... add bonded port PORT in BRIDGE from IFACES del-port [BRIDGE] PORT delete PORT (which may be bonded) from BRIDGE port-to-br PORT print name of bridge that contains PORT

Interface commands (a bond consists of multiple interfaces): list-ifaces BRIDGE print the names of all interfaces on BRIDGE iface-to-br IFACE print name of bridge that contains IFACE

Controller commands: get-controller BRIDGE print the controllers for BRIDGE del-controller BRIDGE delete the controllers for BRIDGE set-controller BRIDGE TARGET... set the controllers for BRIDGE get-fail-mode BRIDGE print the fail-mode for BRIDGE del-fail-mode BRIDGE delete the fail-mode for BRIDGE set-fail-mode BRIDGE MODE set the fail-mode for BRIDGE to MODE

Manager commands: get-manager print the managers del-manager delete the managers set-manager TARGET... set the list of managers to TARGET...

SSL commands: get-ssl print the SSL configuration del-ssl delete the SSL configuration set-ssl PRIV-KEY CERT CA-CERT set the SSL configuration

Switch commands: emer-reset reset switch to known good state

Database commands: list TBL [REC] list RECord (or all records) in TBL find TBL CONDITION... list records satisfying CONDITION in TBL get TBL REC COL[:KEY] print values of COLumns in RECord in TBL set TBL REC COL[:KEY]=VALUE set COLumn values in RECord in TBL add TBL REC COL [KEY=]VALUE add (KEY=)VALUE to COLumn in RECord in TBL remove TBL REC COL [KEY=]VALUE remove (KEY=)VALUE from COLumn clear TBL REC COL clear values from COLumn in RECord in TBL create TBL COL[:KEY]=VALUE create and initialize new record destroy TBL REC delete RECord from TBL wait-until TBL REC [COL[:KEY]=VALUE] wait until condition is true

Potentially unsafe database commands require –force option.

Options:

--db=DATABASE	connect to DATABASE (default: unix:/var/run/openvswitch/db.sock)
--no-wait	do not wait for ovs-vswitchd to reconfigure
--retry	keep trying to connect to server forever
-t, --timeout=SECS	wait at most SECS seconds for ovs-vswitchd
--dry-run	do not commit changes to database
--oneline	print exactly one line of output per command

Logging options:

-vSPEC, --verbose=SPEC	set logging levels
-v, --verbose	set maximum verbosity level

–log-file[=FILE] enable logging to specified FILE (default: /var/log/openvswitch/ovs-vsctl.log)

–syslog-target=HOST:PORT also send syslog msgs to HOST:PORT via UDP –no-syslog equivalent to –verbose=vsctl:syslog:warn

Active database connection methods: tcp:IP:PORT PORT at remote IP ssl:IP:PORT SSL PORT at remote IP unix:FILE Unix domain socket named FILE

Passive database connection methods: ptcp:PORT[:IP] listen to TCP PORT on IP pssl:PORT[:IP] listen for SSL on PORT on IP punix:FILE listen on Unix domain socket FILE

PKI configuration (required to use SSL):

-p, --private-key=FILE	file with private key

-c, --certificate=FILE   file with certificate for private key

-C, --ca-cert=FILE   file with peer CA certificate

Other options:

-h, --help   display this help message

-V, --version   display version information

7.3.14 6.3.14 OpenvSwitch in Allinone

All in one with ens8

ovs-vsctl add-br br-ens8

ovs-vsctl add-port br-ens8 ens8

ifconfig br-ens8 10.3.4.4 up

ip link set br-ens8 promisc on

ip link add proxy-br-eth1 type veth peer name eth1-br-proxy

ip link add proxy-br-ex type veth peer name ex-br-proxy

ovs-vsctl add-br br-eth1

ovs-vsctl add-br br-ex

ovs-vsctl add-port br-eth1 eth1-br-proxy

ovs-vsctl add-port br-ex ex-br-proxy

ovs-vsctl add-port br-ens8 proxy-br-eth1

ovs-vsctl add-port br-ens8 proxy-br-ex

ip link set eth1-br-proxy up promisc on

ip link set ex-br-proxy up promisc on

ip link set proxy-br-eth1 up promisc on

ip link set proxy-br-ex up promisc on

*router ping

ip netns

ip netns exec qdhcp-9cbd5dd0-928a-4808-ae34-4cc2563fa619 ip addr

route add -net 192.168.32.0/24 gw 10.3.4.100

10.3.4.4->10.3.4.100->192.168.32.1 Ok

7.3.15 6.3.15 openstack Allinone

packstack_uninstall.sh

• httpd.service error

mv 10-keystone_wsgi_admin.conf 10-keystone_wsgi_admin.conf.back mv 10-keystone_wsgi_main.conf 10-keystone_wsgi_main.conf.back

and systemctl start httpd.service

7.3.16 6.3.16 openstack Neutron

source keystonerc_osx # neutron net-create osxnet # neutron subnet-create osxnet 192.168.32.0/24 –name osx_subnet –dns-nameserver 8.8.8.8 # neutron net-create ext_net –router:external=True

neutron subnet-create –gateway 10.3.4.1 –disable-dhcp –allocation-pool start=10.3.4.100,end=10.3.4.200 ext_net 10.3.4.0/24 –name ext_subnet neutron subnet-create –disable-dhcp –allocation-pool start=10.3.4.100,end=10.3.4.200 ext_net 10.3.4.0/24 –name ext_subnet

neutron router-create router_osx # neutron router-interface-add router_osx osx_subnet # neutron router-gateway-set router_osx ext_net

vi /root/allinone-answers.cfg

CONFIG_NEUTRON_OVS_VLAN_RANGES=physnet1:10:20 CONFIG_NEUTRON_OVS_BRIDGE_MAPPINGS=physnet1:br-ex

vi /etc/sysconfig/network-scripts/ifcfg-br-ex DEVICE=br-ex DEVICETYPE=ovs TYPE=OVSBridge BOOT-PROTO=none IPADDR=10.20.0.20 NETMASK=255.255.255.0 GATEWAY=10.20.0.1 DNS1=8.8.8.8 DNS2=8.8.4.4 ONBOOT=yes

vi /etc/sysconfig/network-scripts/ifcfg-eth0 DEVICE=eth0 TYPE=OVSPort DEVICETYPE=ovs OVS_BRIDGE=br-ex NM_CONTROLLED=no ONBOOT=yes IPV6INIT=no USERCTL=no

vi /etc/neutron/l3_agent.ini external_network_bridge = br-ens8

ip link set br-ens8 promisc on

- router iptables problem

ip netns ip netns exec qrouter-742cd9c5-de1d-409e-a138-e120f2658222 iptables -S -t nat ip netns exec qrouter-742cd9c5-de1d-409e-a138-e120f2658222 vi /etc/sysconfig/iptables

add security rule all icmp,tcp,udp,ssh for default rule * key point ip link set br-ens8 promisc on

ip netns exec qrouter-f39e7f50-5113-414c-98fa-a94dd7976c57 ifconfig ip netns exec qrouter-f39e7f50-5113-414c-98fa-a94dd7976c57 ip link set qg-6b9a9a40-d7 promisc on ip netns exec qrouter-f39e7f50-5113-414c-98fa-a94dd7976c57 ip link set qg-6b9a9a40-d7 promisc on

*DVR (Distributed Virtual Router) Before Juno, when we deploy Openstack in production, there always is a painful point about L3 Agent: High availability and performance bottleneck

7.3.17 6.3.17 openstack Cinder

openstack cinder does not work in box, it need physical volume

*tgt yum install scsi-target-utils

vi /etc/tgt/targets.conf

include /etc/cinder/volumes/*

vi /etc/cinder/cinder.conf enabled_backends=lvmdriver-1,lvmdriver-2

[lvmdriver-1] volume_group=cinder-volumes-1 volume_driver=cinder.volume.drivers.lvm.LVMISCSIDriver volume_backend_name=LVM_iSCSI1

[lvmdriver-2] volume_group=cinder-volumes-2 volume_driver=cinder.volume.drivers.lvm.LVMISCSIDriver volume_backend_name=LVM_iSCSI2

$ cinder type-create lvm1 cinder type-create lvm2 $ cinder type-key lvm1 set volume_backend_name=LVM_iSCSI1 cinder type-key lvm2 set volume_backend_name=LVM_iSCSI2 $ cinder extra-specs-list (just to check the settings are there)

systemctl enable tgtd.service systemctl start tgtd.service

Define an iscsi target name tgtadm –lld iscsi –op new –mode target –tid 1 -T iqn.2015-07.10.3.0.104:storage.disk1

tgtadm –lld iscsi –op show –mode target

tgtadm –lld iscsi –op new –mode logicalunit –tid 1 –lun 1 -b /dev/vdb

acl setting tgtadm –lld iscsi –mode target –op bind –tid 1 -I ALL // ALL tgtadm –lld iscsi –mode target –op bind –tid 1 -I 192.168.2.48 //for special ip tgtadm –lld iscsi –mode target –op bind –tid 1 -I 10.3.0.0/24 // area

tgtadm –lld iscsi –op new –mode target –tid 2 -T iqn.2015-07.10.3.0.104:storage.disk2 tgtadm –lld iscsi –op new –mode logicalunit –tid 2 –lun 1 -b /dev/vdc

tgtadm –lld iscsi –mode account –op new –user ''tom" –password ''tom'`

*file disk dd if=/dev/zero of=/fs.iscsi.disk bs=1M count=512 tgtadm –lld iscsi –op new –mode logicalunit –tid 0 –lun 1 -b /fs.iscsi.disk

tgtadm –lld iscsi –mode target –op show

netstat -tulpn | grep 3260

iscsiadm –mode discovery –type sendtargets –portal 10.3.0.104 not working properly *iscsi initiator

> [root@www ~]# yum -y install iscsi-initiator-utils

[root@www ~]# vi /etc/iscsi/initiatorname.iscsi # change to the same IQN you set on the iSCSI target server

InitiatorName=iqn.2014-12.world.server:www.server.world [root@www ~]# vi /etc/iscsi/iscsid.conf # line 54: un-comment

node.session.auth.authmethod = CHAP # line 58,59: uncomment and specify the username and password you set on the iSCSI target server

node.session.auth.username = username

node.session.auth.password = password [root@www ~]# systemctl start iscsid

[root@www ~]# systemctl enable iscsid # discover target

[root@www ~]# iscsiadm -m discovery -t sendtargets -p 10.3.0.104

10.0.0.30:3260,1 iqn.2014-12.world.server:storage.target00

confirm status after discovery

[root@www ~]# iscsiadm -m node -o show

BEGIN RECORD 6.2.0.873-24 node.name = iqn.2014-12.world.server:storage.target00 node.tpgt = 1 node.startup = automatic node.leading_login = No ... ... node.conn[0].iscsi.IFMarker = No node.conn[0].iscsi.OFMarker = No # END RECORD

login to the target

[root@www ~]# iscsiadm -m node –login

Logging in to [iface: default, target: iqn.2014-12.world.server:storage.target00, portal: 10.0.0.30,3260] (multiple) Login to [iface: default, target: iqn.2014-12.world.server:storage.target00, portal: 10.0.0.30,3260] successful.

confirm the established session

[root@www ~]# iscsiadm -m session -o show

tcp: [1] 10.0.0.30:3260,1 iqn.2014-12.world.server:storage.target00 (non-flash) # confirm the partitions

[root@www ~]# cat /proc/partitions

major minor #blocks name

> **11 0 1999872 sr0** 8 0 157286400 sda 8 1 512000 sda1 8 2 156773376 sda2

253 0 52428800 dm-0 253 1 6225920 dm-1 253 2 98050048 dm-2

> 8 16 20971520 sdb

***

vi /etc/cinder/cinder.conf enabled_backends=lvmdriver-1

[lvmdriver-1] volume_group=cinder-volumes-1 volume_driver=cinder.volume.drivers.lvm.LVMISCSIDriver volume_backend_name=LVM_iSCSI1

pvcreate /dev/vdb pvcreate /dev/sda

vgcreate cinder-volumes-1 /dev/vdb vgcreate cinder-volumes-2 /dev/sda

systemctl restart openstack-cinder-volume.service

$ cinder type-create lvm1 $ cinder type-key lvm1 set volume_backend_name=LVM_iSCSI1

$ cinder type-create lvm_vdb $ cinder type-key lvm_vdb set volume_backend_name=lvm_vdb

$ cinder type-create lvm_sda $ cinder type-key lvm_sda set volume_backend_name=lvm_sda

systemctl restart openstack-cinder-api.service openstack-cinder-backup.service openstack-cinder-scheduler.service openstack-cinder-volume.service

cinder type-list cinder extra-specs-list

7.3.18 6.3.17 openstack Cinder with Glusterfs

http://www.unixmen.com/install-glusterfs-server-client-centos-7/ http://slidedeck.io/jbernard/cinder-configuration

- On controller

yum install glusterfs-fuse

vi /etc/cinder/cinder.conf enabled_backends=cindervol1,cindervol2

[cindervol1] volume_backend_name=GLUSTER1 volume_driver=cinder.volume.drivers.glusterfs.GlusterfsDriver glusterfs_shares_config=/etc/cinder/shares.conf glusterfs_mount_point_base=/var/lib/cinder/mnt/gluster1

[cindervol2] volume_backend_name=GLUSTER2 volume_driver=cinder.volume.drivers.glusterfs.GlusterfsDriver glusterfs_shares_config=/etc/cinder/shares.conf glusterfs_mount_point_base=/var/lib/cinder/mnt/gluster2

$ cinder type-create gfsvol1 $ cinder type-key gfsvol1 set volume_backend_name=GLUSTER1 $ cinder type-create gfsvol2 $ cinder type-key gfsvol2 set volume_backend_name=GLUSTER2

$ cinder extra-specs-list (just to check the settings are there)

$ cinder type-create lvm $ cinder type-key lvm set volume_backend_name=LVM_iSCSI $ cinder extra-specs-list (just to check the settings are there)

vi /etc/cinder/shares.conf

OpenStackServer3:cindervol1 OpenStackServer3:cindervol2

- Gluster Host

gluster peer probe OpenStackServer1 gluster peer probe OpenStackServer3

gluster pool list

>gluster volume create cindervol1 rep 2 transport tcp OpenStackServer3:/var/lib/cinder/volumes OpenStackServer1:/var/lib/cinder/cindervol1 force volume start cindervol1

---

volume create cindervol2 rep 2 transport tcp OpenStackServer3:/var/lib/cinder/volumes2 OpenStack-Server1:/var/lib/cinder/cindervol2 force volume start cindervol2

Create mount point and mount the volume on both nodes:

[root@glusterfs1 ~]# mount -t glusterfs OpenStackServer3:/cindervol1 /var/lib/cinder/mnt/gluster1/

[root@glusterfs2 ~]# mount -t glusterfs OpenStackServer3:/cindervol1 /var/lib/cinder/mnt/gluster1/

systemctl restart openstack-cinder-volume.service

test cinder create –display-name test 2 cinder create –display-name test2 2

7.3.19 6.3.18 openstack Cinder with cindervolumes

create new

[DEFAULT] state_path=/var/lib/cinder api_paste_config=api-paste.ini enable_v1_api=true os-api_volume_listen=0.0.0.0 osapi_volume_listen_port=8776 rootwrap_config=/etc/cinder/rootwrap.conf auth_strategy=keystone # specify Glance server

glance_host=10.3.0.102 glance_port=9292 # specify RabbitMQ server

rabbit_host=10.3.0.102 rabbit_port=5672 # RabbitMQ user for auth

#rabbit_userid=guest rabbit_userid=guest

RabbitMQ user's password for auth

rabbit_password=guest rpc_backend=rabbit # specify iSCSI target (it's just the own IP)

iscsi_ip_address=10.3.0.104 iscsi_port=3260 iscsi_helper=tgtadm scheduler_driver=cinder.scheduler.filter_scheduler.FilterScheduler volume_manager=cinder.volume.manager.VolumeManager volume_api_class=cinder.volume.api.API volumes_dir=$state_path/volumes # auth info for MariaDB

[database] connection=mysql://cinder:password@10.3.0.102/cinder # auth info for Keystone

[keystone_authtoken] auth_host=10.3.0.102 auth_port=35357 auth_protocol=http admin_user=cinder #admin_password=servicepassword admin_password= admin_tenant_name=service

7.3.20 6.3.19 openstack error

Instance failed to spawn : you must call 'aug-init' first to initialize Augeas

out of physical memory

7.4 6.4 OpenStack Juno +OpenDaylight Helium

https://www.rdoproject.org/Helium_OpenDaylight_Juno_OpenStack https://wiki.opendaylight.org/view/OVSDB:Helium_and_Openstack

opendaylight litium

https://wiki.opendaylight.org/view/OpenDaylight_DLUX:DLUX_Karaf_Feature

chapter 7 :grafana

8.1 7.1 Basic install

8.1.1 7.1.1 influxdb+grafana

https://gist.github.com/ashrithr/9224450

wget http://dl.fedoraproject.org/pub/epel/6/x86_64/epel-release-6-8.noarch.rpm rpm -Uvh epel-release-6*.rpm

```
yum install pycairo Django14 python-ldap python-memcached python-sqlite2 bitmap bitmap-fonts-compat \
python-devel python-crypto pyOpenSSL gcc python-zope-filesystem python-zope-interface git gcc-c++ \
zlib-static MySQL-python python-txamqp python-setuptools python-psycopg2 mod_wsgi
```

. *instal grafana with rpm http://docs.grafana.org/installation/rpm/

Install from package file

You can install Grafana using Yum directly.

$ sudo yum install https://grafanarel.s3.amazonaws.com/builds/grafana-2.1.3-1.x86_64.rpm

Or install manually using rpm.

$ sudo yum install initscripts fontconfig $ sudo rpm -Uvh grafana-2.1.3-1.x86_64.rpm

Install via YUM Repository

Add the following to a new file at /etc/yum.repos.d/grafana.repo

[grafana] name=grafana baseurl=https://packagecloud.io/grafana/stable/el/6/$basearch repo_gpgcheck=1 enabled=1 gpgcheck=1 gpgkey=https://packagecloud.io/gpg.key https://grafanarel.s3.amazonaws.com/RPM-GPG-KEY-grafana sslverify=1 sslcacert=/etc/pki/tls/certs/ca-bundle.crt

There is also a testing repository if you want beta or release candidates.

baseurl=https://packagecloud.io/grafana/testing/el/6/$basearch

Then install Grafana via the yum command.

$ sudo yum install grafana

RPM GPG Key

The RPMs are signed, you can verify the signature with this public GPG key. Package details

Installs binary to /usr/sbin/grafana-server Copies init.d script to /etc/init.d/grafana-server Installs default file (environment vars) to /etc/sysconfig/grafana-server Copies configuration file to /etc/grafana/grafana.ini Installs systemd service (if systemd is available) name grafana-server.service The

default configuration uses a log file at /var/log/grafana/grafana.log The default configuration specifies an sqlite3 database at /var/lib/grafana/grafana.db

Start the server (init.d service)

You can start Grafana by running:

$ sudo service grafana-server start

This will start the grafana-server process as the grafana user, which is created during package installation. The default HTTP port is 3000, and default user and group is admin.

To configure the Grafana server to start at boot time:

$ sudo /sbin/chkconfig –add grafana-server

Start the server (via systemd)

$ systemctl daemon-reload $ systemctl start grafana-server $ systemctl status grafana-server

Enable the systemd service to start at boot

sudo systemctl enable grafana-server.service

Environment file

The systemd service file and init.d script both use the file located at /etc/sysconfig/grafana-server for environment variables used when starting the back-end. Here you can override log directory, data directory and other variables. Logging

By default Grafana will log to /var/log/grafana Database

The default configuration specifies a sqlite3 database located at /var/lib/grafana/grafana.db. Please backup this database before upgrades. You can also use MySQL or Postgres as the Grafana database, as detailed on the configuration page. Configuration

The configuration file is located at /etc/grafana/grafana.ini. Go the Configuration page for details on all those options. Adding data sources

> Graphite InfluxDB OpenTSDB

- install InfluxDB

wget http://get.influxdb.org.s3.amazonaws.com/influxdb-0.8.9-1.x86_64.rpm wget http://influxdb.s3.amazonaws.com/influxdb-0.9.2-1.x86_64.rpm

sudo yum localinstall influxdb-0.8.9-1.x86_64.rpm

sudo /etc/init.d/influxdb start

8.2 7.2 logstash forwarder

logstash forwarder + logstash + elasticsearch+ kibana

logstash forwarder + logstash +graphite +grafana

graphite = Carbon cache+whisper+graphite web

stagemonitor + graphite+grafana

8.2.1 7.2.2 logstach forwarder

```
git clone git://github.com/elasticsearch/logstash-forwarder.git
cd logstash-forwarder
go build -o logstash-forwarder
```

. *centos go language setup in epel

```
yum install golang
```

. Packaging it (optional)

```
gem install bundler
bundle install
```

. * gem

```
yum install ruby
yum install rubygems
```

. * install ruby 1.9.3 http://tecadmin.net/install-ruby-1-9-3-or-multiple-ruby-verson-on-centos-6-3-using-rvm/

yum install rpm-build

```
make rpm
```

.

8.2.2 7.2.3 logstach forwarder

test

8.2.3 7.2.4 sta

test

8.3 7.3 ELK

8.3.1 7.3.1 ELK on CentOS7

centos 7 https://www.digitalocean.com/community/tutorials/how-to-install-elasticsearch-logstash-and-kibana-4-on-centos-7

centos 6 https://gist.github.com/ashrithr/c5c03950ef631ac63c43

8.3.2 7.3.2 scullxbones/docker_grafana_statsd_elk

https://github.com/scullxbones/docker_grafana_statsd_elk